

*Gabe*<sup>3</sup>

GabeL@virginia.edu

<https://gabriel-vzh2vs.github.io/SYS3062-Website/>

*UVA's Systems and Information Engineering*

Februry 24th, 2026

# SYS 3062: Lab 5

*Uncertainty and Output Analysis: The Game*

## The Source Material

Based on Eliyahu Goldratt's novel *The Goal* (Theory of Constraints).

- › **Scenario:** A group of Boy Scouts hiking (a manufacturing line).
- › **The Problem:** The line moves only as fast as the slowest hiker (bottleneck) + random fluctuations.

**The Simulation:** A manufacturing line with  $N$  workers.

- › **Process Time:** Random (originally a die roll, here Gamma distributed).
- › **Buffers:** Finite storage space between workers ( $b$ ).
- › **Objective:** Maximize Throughput ( $y$ ) by tuning the system parameters.

# The Dice Game: Physical vs. Digital

In *The Goal*, Alex Rogo uses a line of Boy Scouts to demonstrate manufacturing physics.

## The Physical Game

- **Workers:** The Scouts.
- **Product:** Matches moved from bowl to bowl.
- **Process Time:** Rolling a Die (1 to 6).
- **Dependency:** You cannot move matches you don't have.

## The Digital Model

- **Workers:** simply.Process.
- **Product:** Integer tokens.
- **Process Time:** Gamma Distribution (Continuous version of a die).
- **Dependency:** yield buffer.get().

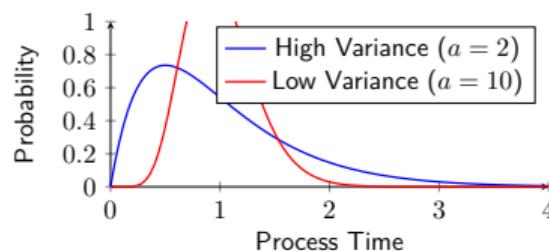
# Modeling Process Time: Die vs. Gamma

## The Uniform Distribution (The Die):

- Mean = 3.5. Variance is fixed.
- *Problem:* Real workers don't have hard "min/max" caps like a die.

## The Gamma Distribution ( $\Gamma$ ):

- We normalize the Mean to 1.0.
- We control Variance using shape parameter  $a$ .
- **Natural Boundaries:** Time cannot be negative, but can occasionally be very long (tail events).



## SimPy Strategy: Blocking & Starving

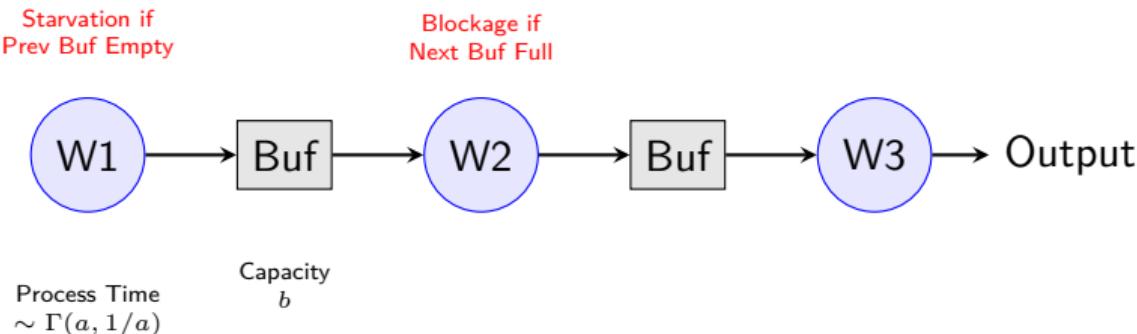
SimPy handles the logic of Blocking (Output Full) and Starving (Input Empty) automatically via 'yield input.get()' and 'yield output.put()'.

```
def worker_process(env, name, input_buffer,
output_buffer, shape_a):
    while True:
        # Wait until item is available in input buffer
        item = yield input_buffer.get()
        time_to_process = np.random.gamma(shape_a, 1.0 /
shape_a)
        yield env.timeout(time_to_process)

        # Wait until space is available in output buffer
        yield output_buffer.put(item)
```

# The Assembly Line Model

A linear sequence of Workers separated by Finite Buffers.



## Setup Step: The Kanban System

**Implementation Strategy:** In SimPy, we use 'simpy.Store(capacity=b)'.

```
def setup_line(env, n, b):
    buffers = []
    # 1. Source (Infinite)
    buffers.append(simpy.Store(env, capacity=inf))
    # 2. Intermediate Buffers (The Kanban Limit)
    for i in range(n - 1):
        buffers.append(simpy.Store(env, capacity=b))
    # 3. Sink (Infinite)
    buffers.append(simpy.Store(env, capacity=inf))
    return buffers
```

*Note: If Worker 1 tries to put an item into Buffer 2, and Buffer 2 has 'b' items, Worker 1 freezes (Blocking). This is how Kanban limits overproduction.*

In the Dice Game, raw speed doesn't matter. Only finished goods matter.

## Throughput Calculation

$$y = \frac{\text{Total Items in Sink}}{\text{Total Simulation Time}}$$

### Setup Requirements:

- 1 **Warm-up:** Pre-fill buffers with 2 items to avoid empty pipe startup bias.
- 2 **Duration:** Run for 1,000 time units to allow statistical fluctuations to average out.
- 3 **Seeding:** Use “`np.random.seed(42)`” to ensure your dice rolls are reproducible.

Now, it's your turn to implement this Queuing System in Python or using Excel! Good Luck!

# How do we analyze a system: Experimental Design (DOE)

## The Problem

A simulation is a "Black Box." We put inputs in, and get an output out. But we don't inherently know *which* input caused the output change.

**The Solution:** experimental design, specifically the sub-method of factorial design  
We systematically vary the input parameters (Factors) in a grid.

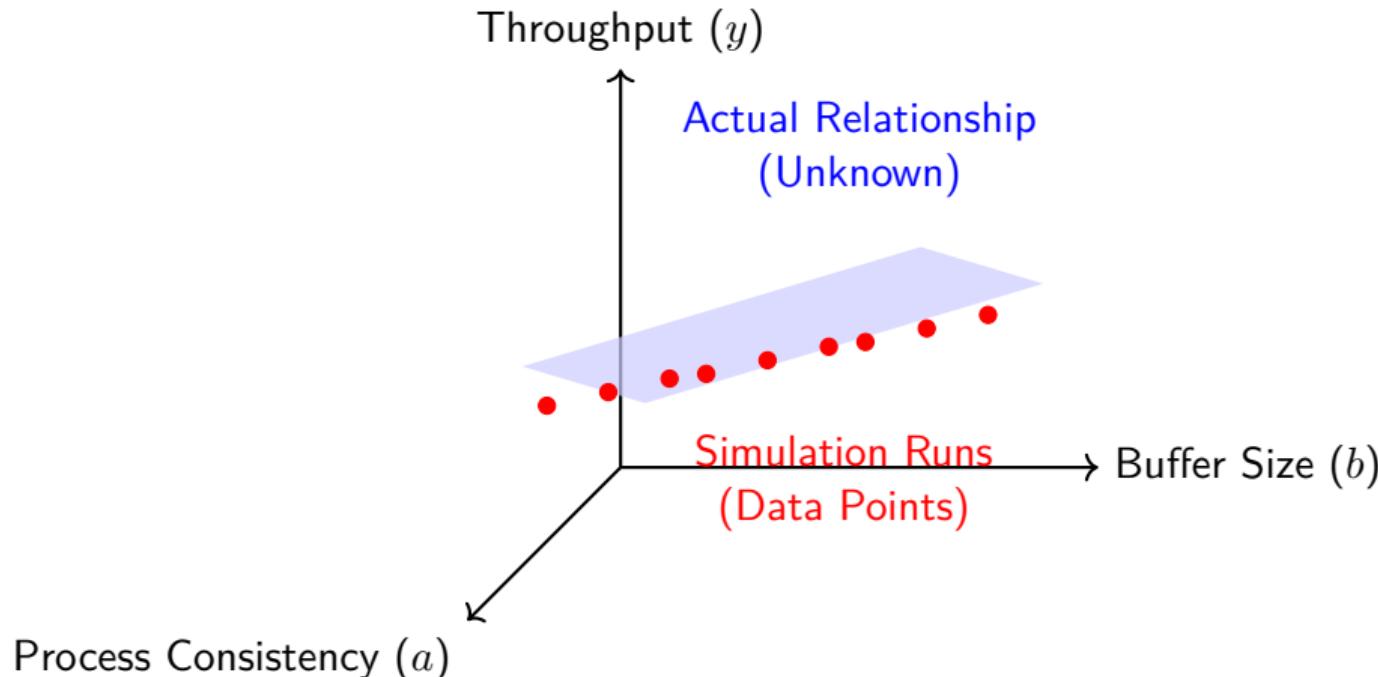
## Lab 6 Factors:

- **$n$  (Line Length):** 11 levels (5, 6, ..., 20).
- **$b$  (Buffer Size):** 10 levels (1, ..., 10).
- **$a$  (Variance):** 6 levels (2, 3, ..., 20).
- **$\sigma$  (Worker Skill):** 5 levels (0...0.1).

$$\text{Total Runs} = 11 \times 10 \times 6 \times 5 = \mathbf{3,300} \text{ Simulations}$$

# Visualizing the Design Space

Instead of guessing random points, we create a structured grid (Hypercube) of scenarios.



# Experimental Design: Choosing the Points

How do we choose which input combinations to run?

## Full Factorial Design ( $L^k$ )

Test **every possible combination** of all factors.

- **Pros:** Captures detailed behavior and complex interactions.
- **Cons:** Curse of Dimensionality. ( $10 \text{ levels}^4 \text{ factors} = 10,000 \text{ runs}$ ).

## $2^k$ Factorial Design

Test only the **Low (-1)** and **High (+1)** levels of each factor.

- **Pros:** Extremely efficient ( $2^4 = 16 \text{ runs}$ ). Great for "Screening" (finding what matters).
- **Cons:** Assumes linearity. Cannot see curvature (peaks/valleys) inside the range.

What if  $2^k$  is still too expensive? (e.g., 20 factors  $\rightarrow$  1,000,000 runs).

### 3. Fractional Factorial ( $2^{k-p}$ )

We run only a structured subset (e.g., a Half-Fraction).

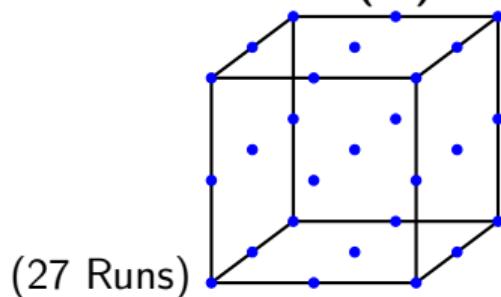
- **Logic:** "Sparsity of Effects." High-order interactions (e.g., Factor A  $\times$  B  $\times$  C  $\times$  D) are usually negligible.
- **Trade-off:** Aliasing. We might confuse Main Effects with Interaction Effects.

*Example: In a  $2^{3-1}$  design, we run 4 experiments instead of 8. We lose some detail but save 50% of the computing cost.*

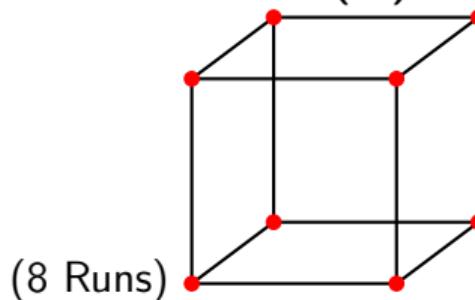
# Visualizing Factorial Spaces

Comparing a 3-Factor Experiment ( $k = 3$ ).

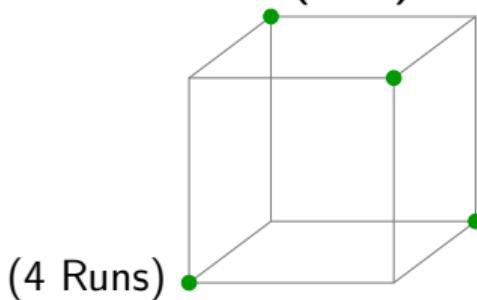
**Full Factorial** ( $3^3$ )



**$2^k$  Factorial** ( $2^3$ )



**Fractional** ( $2^{3-1}$ )



We run a **Full Factorial Experiment** (3,300 runs) to understand how four parameters affect throughput ( $y$ ).

## Input Factors ( $X$ ):

- **$n$  (Line Length):** 5 to 20 workers.
- **$b$  (Buffer Size):** 1 to 10 items.
- **$a$  (Process Consistency):** Shape parameter (Higher  $a$  = Less Variance).
- **$\sigma$  (Worker Variance):** Skill variation between workers.

## Statistical Model (OLS):

$$y \sim 1 + n + \log(1 + b) \times \log(a) + \sigma$$

We hypothesize an interaction between Buffer Size ( $b$ ) and Process Consistency ( $a$ ).

These conclusions are backed by a *meta-model* on top of the full-factorial simulation.

## Key Insights:

- **Buffers Help:** Increasing buffer size ( $b$ ) significantly improves throughput, but with diminishing returns (Logarithmic).
- **Variance Kills:** High process variance (Low  $a$ ) destroys throughput, even with buffers.
- **Line Length Hurts:** Longer lines ( $n$ ) generally have lower throughput due to "Dependency Accumulation".

# How can we interpret complex systems: Metamodels?

## Definition

A metamodel is a "Model of a Model."

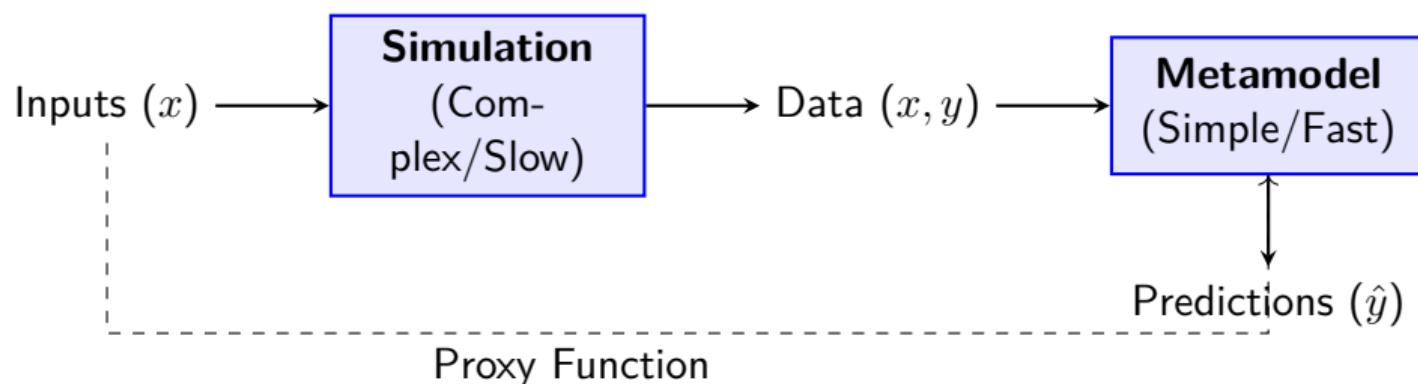
- **The Simulation:** A complex, slow, stochastic representation of reality (e.g., Our SimPy Factory).
- **The Metamodel:** A simple, fast, mathematical approximation of the simulation (e.g., Linear Regression).

## Why use them?

- **Speed:** Running the simulation 1,000 times takes an hour. Evaluating the regression equation takes milliseconds.
- **Optimization:** It is easier to find the maximum of a smooth equation ( $y = \beta x$ ) than a noisy simulation.
- **Insight:** Coefficients reveal *sensitivity* (which factors matter most).

# The Metamodeling Workflow

We do not replace the simulation; we use it to *train* the metamodel.



Once we have the data, we use Ordinary Least Squares (OLS) or a Generalized Linear Model (GLM) to fit a mathematical equation to our simulation results.

## Why OLS?

- 1 **Quantification:** It tells us exactly *how much*  $y$  changes when we increase  $b$  by 1 unit.
- 2 **Significance:** It tells us if a factor matters at all (P-Values).
- 3 **Metamodeling:** The regression equation becomes a Fast Proxy. We can predict throughput without running the slow simulation again.

## The Lab 6 Model:

$$\hat{y} = \beta_0 + \beta_1 n + \beta_2 \log(1 + b) + \beta_3 \log(a) + \dots$$

# Regression Output (Statsmodels)

We use a GLM (in OLS mode here) Regression to quantify the effects.

Factor	Coef	P-Value	Interpretation
Intercept	0.669	0.000	Baseline Throughput
$n$ (Length)	-0.001	0.181	Minimal negative impact
$\log(1 + b)$ (Buffer)	<b>0.105</b>	<b>0.000</b>	<b>Strong Positive Effect</b>
$\log(a)$ (Consistency)	<b>0.106</b>	<b>0.000</b>	<b>Strong Positive Effect</b>
Interaction $b : a$	-0.035	0.008	Buffers matter less if variance is low
$\sigma$ (Worker Var)	-0.749	0.000	Strong Negative Effect

## Conclusion

To fix the factory: **Reduce Variance first, then Add Buffers.**

# Interpreting the Regression Table

How to read the statsmodels output from Lab 6:

Component	Example Value	Meaning
<b>Coefficient</b>	+0.105 ( $\log b$ )	<b>Magnitude.</b> Increasing buffer size raises throughput.
	-0.749 ( $\sigma$ )	Increasing worker variance <i>hurts</i> throughput.
<b>P-Value</b>	0.000	<b>Certainty.</b> This effect is real, not random noise.
	0.181 ( $n$ )	This factor (Line Length) might not matter statistically.
<b>R-Squared</b>	0.735	<b>Fit.</b> Our model explains 73.5% of the system's behavior.

## The Interaction Term

We found a significant interaction:  $\log(1 + b) \times \log(a)$ . *Translation: Buffers are extremely useful when variance is high (low  $a$ ), but less useful when the process is already stable.*

Throughout the course, I may ask for your feedback on:

- Course Materials
- Lecture & Lab
- Assignments



First, I will thank you all again for coming to lab at 5 pm on a Tuesday!

And now, you need to do the following tasks:

- **Submit** your Python file to Gradescope by Spring Break!
- **Have** fun on spring break!

Your exams will be graded sooner than you think!™